

```
In [1]: library(boot)
library(glmnet)
```

```
Loading required package: Matrix
Loading required package: foreach
Loaded glmnet 2.0-16
```

```
In [2]: set.seed(2019)
```

```
In [3]: train <- read.csv("trainC.csv")
test <- read.csv("testC.csv")
train <- subset(train, select = -c(sessionDate, trialNum, timeSinceKetamine))
test <- subset(test, select = -c(sessionDate, trialNum, timeSinceKetamine))

#TRAIN 1: ALL COVARIATES PLUS INTERACTION TERMS
train1 <- read.csv("trainC.csv")
test1 <- read.csv("testC.csv")
train1 <- subset(train1, select = -c(sessionDate, trialNum, timeSinceKetamine))
test1 <- subset(test1, select = -c(sessionDate, trialNum, timeSinceKetamine))

#TRAIN 2: ALL COVARIATES NO INTERACTION TERMS
train2 <- subset(train1, select = c(totalCellNum,gender,genotype,weight_g,ketamine_day,
correlationScore,lickAccuracy,lickNumber,avgFR,
avgSingleCellVariance,varianceFR,avgTrialSpeed,
varianceSpeed,medianCellDepth,ketBool))
test2 <- subset(test1, select = c(totalCellNum,gender,genotype,weight_g,ketamine_day,
correlationScore,lickAccuracy,lickNumber,avgFR,
avgSingleCellVariance,varianceFR,avgTrialSpeed,
varianceSpeed,medianCellDepth,ketBool))
```

```
In [4]: # First, Let's do a 50% split on the training data to determine the best Lambda
n = length(train[,1])
n50 = round(n/2)
train50A = train[1:n50,]
train50B = train[(n50+1):n,]
```

Basic Logistic Regression Model with Interaction Terms

Estimate test error

```

In [5]: k = 10
n = length(train1[,1])
fsize = round(n/k)
rmse = rep(0,k)
zloss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  df_train <- train1[-(((i-1)*fsize+1):(i*fsize)),]
  df_val <- train1[(((i-1)*fsize+1):(i*fsize)),]
  # Fit model on training and make predictions on validation
  model_cv <- glm(ketBool ~ . + animalName:correlationScore
                  + animalName:lickAccuracy
                  + animalName:lickNumber
                  + animalName:avgFR
                  + animalName:avgSingleCellVariance
                  + animalName:varianceFR
                  + animalName:avgTrialSpeed
                  + animalName:varianceSpeed, data=df_train, family='binomial')
  lr_pred_lo <- predict(model_cv,df_val) # Lo : Log odds
  num_val = length(df_val$ketBool)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (lr_pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = df_val$ketBool[j]
  }
  # Compute 0-1 loss for each observation
  lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 loss on the val set
  zloss[i] = mean(lr_loss)
}
df_train <- train1[-(((k-1)*fsize+1):n),]
df_val <- train1[(((k-1)*fsize+1):n),]
# Fit model on training and make predictions on validation
model_cv <- glm(ketBool ~ . + animalName:correlationScore
                + animalName:lickAccuracy
                + animalName:lickNumber
                + animalName:avgFR
                + animalName:avgSingleCellVariance
                + animalName:varianceFR
                + animalName:avgTrialSpeed
                + animalName:varianceSpeed, data=df_train, family='binomial')
lr_pred_lo <- predict(model_cv,df_val) # Lo : Log odds
num_val = length(df_val$ketBool)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (lr_pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = df_val$ketBool[j]
}
lr_loss = abs(lr_pred-actual)
zloss[k] = mean(lr_loss)
test_error_est = mean(zloss)

cat("=====\n")
cat("Logistic Regression Model with Interaction Terms\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")

```



```
In [6]: summary(model_cv)
```

```
Call:
glm(formula = ketBool ~ . + animalName:correlationScore + animalName:lickAccuracy +
  animalName:lickNumber + animalName:avgFR + animalName:avgSingleCellVariance +
  animalName:varianceFR + animalName:avgTrialSpeed + animalName:varianceSpeed,
  family = "binomial", data = df_train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.5384	-0.1027	0.0000	0.0712	5.5159

Coefficients: (18 not defined because of singularities)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-27.2669	6.0252	-4.526	6.03e-06 ***
animalNameG2	0.4197	3.3163	0.127	0.899281
animalNameG3	-0.9364	3.5589	-0.263	0.792468
animalNameG4	1.4332	3.8333	0.374	0.708496
animalNameG5	-8.9931	4.1376	-2.173	0.029743 *
animalNameHCN1	-14.0175	10.5442	-1.329	0.183715
animalNameHCNb2	8.4330	4.4879	1.879	0.060239 .
animalNameHCNb4	-4.7334	23.8994	-0.198	0.843002
animalNameHCNd1	104.4537	16.9813	6.151	7.69e-10 ***
animalNameHCNd2	68.3925	13.2492	5.162	2.44e-07 ***
animalNameHCNe1	57.9201	13.7486	4.213	2.52e-05 ***
animalNameHCNe2	102.2463	16.6521	6.140	8.25e-10 ***
animalNameHCNe3	66.1861	12.7951	5.173	2.31e-07 ***
animalNamepI1	-7.5109	6.2653	-1.199	0.230603
totalCellNum	1.4313	1.2711	1.126	0.260152
gender	NA	NA	NA	NA
genotype	NA	NA	NA	NA
weight_g	-42.4023	7.1213	-5.954	2.61e-09 ***
ketamine_day	0.8366	0.7416	1.128	0.259265
correlationScore	0.4354	7.9683	0.055	0.956425
lickAccuracy	-9.0325	5.2510	-1.720	0.085409 .
lickNumber	16.5052	8.4302	1.958	0.050246 .
avgFR	-116.8217	22.1238	-5.280	1.29e-07 ***
avgSingleCellVariance	118.9793	28.1016	4.234	2.30e-05 ***
varianceFR	-11.7773	9.6821	-1.216	0.223831
avgTrialSpeed	-9.3312	5.5556	-1.680	0.093036 .
varianceSpeed	17.2301	6.8436	2.518	0.011813 *
medianCellDepth	4.3650	0.8865	4.924	8.49e-07 ***
totalCellNumxCorrelationScore	-3.1603	1.0375	-3.046	0.002319 **
totalCellNumxLickAccuracy	0.8750	0.5696	1.536	0.124463
totalCellNumxLickNumber	1.6250	0.7325	2.218	0.026530 *
totalCellNumxAvgFR	0.8384	1.7622	0.476	0.634260
totalCellNumxAvgSingleCellVariance	2.1842	1.5135	1.443	0.148980
totalCellNumxVarianceFR	-0.6762	0.5323	-1.270	0.203957
totalCellNumxAvgTrialSpeed	-0.7475	0.9554	-0.782	0.433959
totalCellNumxVarianceSpeed	1.2069	0.6701	1.801	0.071709 .
genderxCorrelationScore	4.7113	4.3436	1.085	0.278070
genderxLickAccuracy	-3.1125	2.0324	-1.531	0.125657
genderxLickNumber	4.7001	2.6640	1.764	0.077679 .
genderxAvgFR	-87.0050	14.6140	-5.954	2.62e-09 ***
genderxAvgSingleCellVariance	40.7485	11.0054	3.703	0.000213 ***
genderxVarianceFR	-4.6695	5.1769	-0.902	0.367069
genderxAvgTrialSpeed	-9.4424	3.8540	-2.450	0.014284 *
genderxVarianceSpeed	4.5012	3.7641	1.196	0.231762
genotypexCorrelationScore	3.0093	2.0973	1.435	0.151338
genotypexLickAccuracy	-0.2219	1.1406	-0.195	0.845758
genotypexLickNumber	4.0348	1.9529	2.066	0.038822 *
genotypexAvgFR	-42.7736	8.1363	-5.257	1.46e-07 ***
genotypexAvgSingleCellVariance	37.2330	8.5842	4.337	1.44e-05 ***
genotypexVarianceFR	-2.6952	2.6845	-1.004	0.315386
genotypexAvgTrialSpeed	-1.6597	2.0224	-0.821	0.411833
genotypexVarianceSpeed	-0.2610	1.7360	-0.150	0.880487
weight_gxCorrelationScore	-8.5846	10.1501	-0.846	0.397685
weight_gxLickAccuracy	10.7017	7.2068	1.485	0.137560
weight_gxLickNumber	-26.7826	11.5589	-2.317	0.020501 *
weight_gxAvgFR	193.5887	31.4637	6.153	7.61e-10 ***
weight_gxAvgSingleCellVariance	-154.5028	30.8680	-5.005	5.58e-07 ***
weight_gxVarianceFR	18.4687	14.8475	1.244	0.213540
weight_gxAvgTrialSpeed	13.0744	6.7444	1.939	0.052555 .
weight_gxVarianceSpeed	-17.5478	8.0375	-2.183	0.029019 *
ketamine_dayxCorrelationScore	2.0090	0.6098	3.295	0.000985 ***
ketamine_dayxLickAccuracy	-0.2291	0.4973	-0.461	0.645009
ketamine_dayxLickNumber	0.6505	0.5670	1.147	0.251278
ketamine_dayxAvgFR	-7.5420	2.7782	-2.715	0.006633 **
ketamine_dayxAvgSingleCellVariance	0.7418	2.6459	0.280	0.779199
ketamine_dayxVarianceFR	-0.9275	1.0399	-0.892	0.372456
ketamine_dayxAvgTrialSpeed	1.3230	0.7267	1.820	0.068686 .
ketamine_dayxVarianceSpeed	0.1314	0.8028	0.164	0.869944
medianCellDepthxCorrelationScore	-1.4019	0.5466	-2.565	0.010321 *

medianCellDepthxLickAccuracy	-0.3065	0.4351	-0.704	0.481208
medianCellDepthxLickNumber	0.4651	0.7274	0.639	0.522605
medianCellDepthxAvgFR	-14.1645	2.7873	-5.082	3.74e-07 ***
medianCellDepthxAvgSingleCellVariance	8.6487	4.3498	1.988	0.046779 *
medianCellDepthxVarianceFR	0.6166	0.7940	0.777	0.437441
medianCellDepthxAvgTrialSpeed	1.5695	0.7607	2.063	0.039087 *
medianCellDepthxVarianceSpeed	-2.1309	0.9237	-2.307	0.021060 *
animalNameG2:correlationScore	1.5773	1.9963	0.790	0.429456
animalNameG3:correlationScore	1.3056	2.4116	0.541	0.588254
animalNameG4:correlationScore	0.8286	2.3078	0.359	0.719573
animalNameG5:correlationScore	1.3590	2.7886	0.487	0.626014
animalNameHCN1:correlationScore	7.7178	4.1183	1.874	0.060928 .
animalNameHCNb2:correlationScore	5.6795	3.0742	1.847	0.064681 .
animalNameHCNb4:correlationScore	14.4818	6.8450	2.116	0.034373 *
animalNameHCNd1:correlationScore	0.9176	2.0826	0.441	0.659487
animalNameHCNd2:correlationScore	4.9989	1.7497	2.857	0.004276 **
animalNameHCNe1:correlationScore	1.6307	3.1039	0.525	0.599324
animalNameHCNe2:correlationScore	NA	NA	NA	NA
animalNameHCNe3:correlationScore	NA	NA	NA	NA
animalNamenpI1:correlationScore	7.3425	2.8767	2.552	0.010697 *
animalNameG2:lickAccuracy	-1.7099	1.1916	-1.435	0.151303
animalNameG3:lickAccuracy	0.4135	0.9017	0.459	0.646531
animalNameG4:lickAccuracy	-0.1212	1.1478	-0.106	0.915904
animalNameG5:lickAccuracy	2.3511	1.7690	1.329	0.183839
animalNameHCN1:lickAccuracy	-2.1156	2.6306	-0.804	0.421274
animalNameHCNb2:lickAccuracy	-3.7346	1.6714	-2.234	0.025455 *
animalNameHCNb4:lickAccuracy	-5.7505	3.6852	-1.560	0.118654
animalNameHCNd1:lickAccuracy	0.5928	1.1011	0.538	0.590326
animalNameHCNd2:lickAccuracy	-3.2750	1.8848	-1.738	0.082279 .
animalNameHCNe1:lickAccuracy	2.7730	2.1648	1.281	0.200209
animalNameHCNe2:lickAccuracy	NA	NA	NA	NA
animalNameHCNe3:lickAccuracy	NA	NA	NA	NA
animalNamenpI1:lickAccuracy	-2.8775	1.1676	-2.464	0.013723 *
animalNameG2:lickNumber	-0.1500	1.0227	-0.147	0.883379
animalNameG3:lickNumber	-1.8930	1.2920	-1.465	0.142882
animalNameG4:lickNumber	-3.7432	1.6083	-2.327	0.019944 *
animalNameG5:lickNumber	-9.2212	2.5582	-3.605	0.000313 ***
animalNameHCN1:lickNumber	9.9468	4.2924	2.317	0.020488 *
animalNameHCNb2:lickNumber	2.1215	3.2220	0.658	0.510251
animalNameHCNb4:lickNumber	4.2728	8.6890	0.492	0.622893
animalNameHCNd1:lickNumber	1.6771	1.2682	1.322	0.186031
animalNameHCNd2:lickNumber	0.9823	1.4131	0.695	0.486990
animalNameHCNe1:lickNumber	-8.0085	3.3984	-2.357	0.018446 *
animalNameHCNe2:lickNumber	NA	NA	NA	NA
animalNameHCNe3:lickNumber	NA	NA	NA	NA
animalNamenpI1:lickNumber	-4.2673	3.6055	-1.184	0.236592
animalNameG2:avgFR	7.4748	3.6790	2.032	0.042177 *
animalNameG3:avgFR	23.6143	5.1125	4.619	3.86e-06 ***
animalNameG4:avgFR	31.2290	4.7689	6.548	5.81e-11 ***
animalNameG5:avgFR	39.7199	6.5348	6.078	1.22e-09 ***
animalNameHCN1:avgFR	-45.6980	12.3054	-3.714	0.000204 ***
animalNameHCNb2:avgFR	-35.6108	7.8888	-4.514	6.36e-06 ***
animalNameHCNb4:avgFR	-9.3360	62.1295	-0.150	0.880554
animalNameHCNd1:avgFR	9.1863	3.9552	2.323	0.020202 *
animalNameHCNd2:avgFR	-13.7740	6.7154	-2.051	0.040256 *
animalNameHCNe1:avgFR	56.3254	8.8355	6.375	1.83e-10 ***
animalNameHCNe2:avgFR	NA	NA	NA	NA
animalNameHCNe3:avgFR	NA	NA	NA	NA
animalNamenpI1:avgFR	-6.9399	8.7446	-0.794	0.427414
animalNameG2:avgSingleCellVariance	-5.8451	7.1199	-0.821	0.411667
animalNameG3:avgSingleCellVariance	-21.4194	8.6198	-2.485	0.012958 *
animalNameG4:avgSingleCellVariance	-21.1395	8.1704	-2.587	0.009673 **
animalNameG5:avgSingleCellVariance	-43.8020	9.4607	-4.630	3.66e-06 ***
animalNameHCN1:avgSingleCellVariance	57.2169	15.6053	3.667	0.000246 ***
animalNameHCNb2:avgSingleCellVariance	26.6612	11.7217	2.275	0.022935 *
animalNameHCNb4:avgSingleCellVariance	53.1048	30.8837	1.720	0.085522 .
animalNameHCNd1:avgSingleCellVariance	-0.8131	6.4811	-0.125	0.900156
animalNameHCNd2:avgSingleCellVariance	30.6919	12.4083	2.473	0.013380 *
animalNameHCNe1:avgSingleCellVariance	-38.3291	10.3293	-3.711	0.000207 ***
animalNameHCNe2:avgSingleCellVariance	NA	NA	NA	NA
animalNameHCNe3:avgSingleCellVariance	NA	NA	NA	NA
animalNamenpI1:avgSingleCellVariance	-12.8571	16.3104	-0.788	0.430536
animalNameG2:varianceFR	0.7852	1.8165	0.432	0.665556
animalNameG3:varianceFR	1.3876	2.6141	0.531	0.595547
animalNameG4:varianceFR	2.3089	2.5615	0.901	0.367378
animalNameG5:varianceFR	4.1829	3.0171	1.386	0.165618
animalNameHCN1:varianceFR	-3.2583	4.8433	-0.673	0.501115
animalNameHCNb2:varianceFR	-3.6846	3.1546	-1.168	0.242803
animalNameHCNb4:varianceFR	-25.8202	31.0535	-0.831	0.405705
animalNameHCNd1:varianceFR	0.1037	1.9293	0.054	0.957136
animalNameHCNd2:varianceFR	-2.8652	2.7017	-1.061	0.288911
animalNameHCNe1:varianceFR	4.8279	4.0581	1.190	0.234157

animalNameHCNe2:varianceFR	NA	NA	NA	NA
animalNameHCNe3:varianceFR	NA	NA	NA	NA
animalNamepI1:varianceFR	4.5827	6.4964	0.705	0.480550
animalNameG2:avgTrialSpeed	-2.4850	1.7024	-1.460	0.144372
animalNameG3:avgTrialSpeed	-1.8267	1.9955	-0.915	0.359969
animalNameG4:avgTrialSpeed	-1.1678	1.9046	-0.613	0.539752
animalNameG5:avgTrialSpeed	-0.8649	2.2433	-0.386	0.699814
animalNameHCN1:avgTrialSpeed	-7.7469	3.0655	-2.527	0.011499 *
animalNameHCNb2:avgTrialSpeed	-6.3337	2.3588	-2.685	0.007251 **
animalNameHCNb4:avgTrialSpeed	-6.3992	4.3391	-1.475	0.140272
animalNameHCNd1:avgTrialSpeed	-0.6755	1.3358	-0.506	0.613098
animalNameHCNd2:avgTrialSpeed	0.7792	1.2862	0.606	0.544644
animalNameHCNe1:avgTrialSpeed	2.9309	2.3740	1.235	0.216977
animalNameHCNe2:avgTrialSpeed	NA	NA	NA	NA
animalNameHCNe3:avgTrialSpeed	NA	NA	NA	NA
animalNamepI1:avgTrialSpeed	-4.7091	2.0490	-2.298	0.021549 *
animalNameG2:varianceSpeed	0.9403	2.2028	0.427	0.669485
animalNameG3:varianceSpeed	-0.8958	2.6235	-0.341	0.732764
animalNameG4:varianceSpeed	-3.0501	2.4720	-1.234	0.217244
animalNameG5:varianceSpeed	-3.5242	2.8684	-1.229	0.219219
animalNameHCN1:varianceSpeed	2.9468	4.3016	0.685	0.493319
animalNameHCNb2:varianceSpeed	3.0695	3.3371	0.920	0.357673
animalNameHCNb4:varianceSpeed	-20.5634	18.9175	-1.087	0.277034
animalNameHCNd1:varianceSpeed	-1.3491	1.9970	-0.676	0.499317
animalNameHCNd2:varianceSpeed	-4.5265	2.2934	-1.974	0.048410 *
animalNameHCNe1:varianceSpeed	-5.4418	2.5092	-2.169	0.030103 *
animalNameHCNe2:varianceSpeed	NA	NA	NA	NA
animalNameHCNe3:varianceSpeed	NA	NA	NA	NA
animalNamepI1:varianceSpeed	-3.9996	2.8047	-1.426	0.153858

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 4989.22 on 3599 degrees of freedom
Residual deviance: 876.32 on 3438 degrees of freedom
AIC: 1200.3

Number of Fisher Scoring iterations: 12

Reduced dataset to match Lasso and Ridge

Basic Logistic Regression without Interaction Terms

```
In [8]: k = 10
n = length(train2[,1])
fsize = round(n/k)
rmse = rep(0,k)
zoloss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  df_train <- train2[-(((i-1)*fsize+1):(i*fsize)),]
  df_val <- train2[((i-1)*fsize+1):(i*fsize),]
  # Fit model on training and make predictions on validation
  model_cv <- glm(ketBool ~ ., data=df_train, family='binomial')
  lr_pred_lo <- predict(model_cv,df_val) # Lo : Log odds
  num_val = length(df_val$ketBool)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (lr_pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = df_val$ketBool[j]
  }
  # Compute 0-1 Loss for each observation
  lr_loss = abs(lr_pred-actual) # Loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 Loss on the val set
  zoloss[i] = mean(lr_loss)
}
df_train <- train2[-(((k-1)*fsize+1):n),]
df_val <- train2[((k-1)*fsize+1):n,]
# Fit model on training and make predictions on validation
model_cv <- glm(ketBool ~ ., data=df_train, family='binomial')
lr_pred_lo <- predict(model_cv,df_val) # Lo : Log odds
num_val = length(df_val$ketBool)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (lr_pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = df_val$ketBool[j]
}
lr_loss = abs(lr_pred-actual)
zoloss[k] = mean(lr_loss)
test_error_est = mean(zoloss)

cat("=====\n")
cat("Logistic Regression Model without Interaction Terms\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")
```

=====
Logistic Regression Model without Interaction Terms

Zero-One Loss (10-fold Cross-Validation Average): 0.1413709
Accuracy (10-fold Cross-Validation Average): 0.8586291
=====

GLMNET

```

In [9]: # First, let's do a 50% split on the training data to determine the best lambda
n = length(train[,1])
n50 = round(n/2)
train50A = train[1:n50,]
train50B = train[(n50+1):n,]

xA = model.matrix(ketBool ~ . + animalName:correlationScore
                  + animalName:lickAccuracy
                  + animalName:lickNumber
                  + animalName:avgFR
                  + animalName:avgSingleCellVariance
                  + animalName:varianceFR
                  + animalName:avgTrialSpeed
                  + animalName:varianceSpeed, data = train50A)
yA = train50A$ketBool

xB = model.matrix(ketBool ~ . + animalName:correlationScore
                  + animalName:lickAccuracy
                  + animalName:lickNumber
                  + animalName:avgFR
                  + animalName:avgSingleCellVariance
                  + animalName:varianceFR
                  + animalName:avgTrialSpeed
                  + animalName:varianceSpeed, data = train50B)
yB = train50B$ketBool

```

Lasso

```

In [10]: # Select regularization parameter over trainA (50% of training data)
model_lasso <- cv.glmnet(xA, yA, family='binomial',alpha=1)
lambda_min = model_lasso$lambda.min
lambda_1se = model_lasso$lambda.1se

```

lambda.min

```

In [11]: k = 10
n = length(train50B[,1])
fsize = round(n/k)
rmse = rep(0,k)
zolooss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  xB_train = xB[-(((i-1)*fsize+1):(i*fsize)),]
  yB_train = yB[-(((i-1)*fsize+1):(i*fsize))]
  xB_val = xB[((i-1)*fsize+1):(i*fsize),]
  yB_val = yB[((i-1)*fsize+1):(i*fsize)]
  # Fit model on training and make predictions on validation
  model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=1,lambda=lambda_min)
  pred_lo = predict(model_cv, newx = xB_val)
  num_val = length(yB_val)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = yB_val[j]
  }
  # Compute 0-1 loss for each observation
  lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 loss on the val set
  zolooss[i] = mean(lr_loss)
}
xB_train = xB[-(((k-1)*fsize+1):(length(yB))),]
yB_train = yB[-(((k-1)*fsize+1):(length(yB)))]
xB_val = xB[((k-1)*fsize+1):(length(yB)),]
yB_val = yB[((k-1)*fsize+1):(length(yB))]
# Fit model on training and make predictions on validation
model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=1,lambda=lambda_min)
pred_lo = predict(model_cv, newx = xB_val)
num_val = length(yB_val)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = yB_val[j]
}
# Compute 0-1 loss for each observation
lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
# Compute mean 0-1 loss on the val set
zolooss[k] = mean(lr_loss)
test_error_est = mean(zolooss)

cat("=====\n")
cat("GLMNET Lasso Logistic Regression Model with lambda.min\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")

```

```

=====
GLMNET Lasso Logistic Regression Model with lambda.min

```

```

Zero-One Loss (10-fold Cross-Validation Average): 0.07002764
Accuracy (10-fold Cross-Validation Average): 0.9299724
=====

```

lambda.1se

```

In [12]: k = 10
n = length(train50B[,1])
fsize = round(n/k)
rmse = rep(0,k)
zolooss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  xB_train = xB[-(((i-1)*fsize+1):(i*fsize)),]
  yB_train = yB[-(((i-1)*fsize+1):(i*fsize))]
  xB_val = xB[((i-1)*fsize+1):(i*fsize),]
  yB_val = yB[((i-1)*fsize+1):(i*fsize)]
  # Fit model on training and make predictions on validation
  model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=1,lambda=lambda_1se)
  pred_lo = predict(model_cv, newx = xB_val)
  num_val = length(yB_val)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = yB_val[j]
  }
  # Compute 0-1 loss for each observation
  lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 loss on the val set
  zolooss[i] = mean(lr_loss)
}
xB_train = xB[-(((k-1)*fsize+1):(length(yB))),]
yB_train = yB[-(((k-1)*fsize+1):(length(yB)))]
xB_val = xB[((k-1)*fsize+1):(length(yB)),]
yB_val = yB[((k-1)*fsize+1):(length(yB))]
# Fit model on training and make predictions on validation
model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=1,lambda=lambda_1se)
pred_lo = predict(model_cv, newx = xB_val)
num_val = length(yB_val)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = yB_val[j]
}
# Compute 0-1 loss for each observation
lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
# Compute mean 0-1 loss on the val set
zolooss[k] = mean(lr_loss)
test_error_est = mean(zolooss)

cat("=====\n")
cat("GLMNET Lasso Logistic Regression Model with lambda.1se\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")

```

```

=====
GLMNET Lasso Logistic Regression Model with lambda.1se

Zero-One Loss (10-fold Cross-Validation Average): 0.07403518
Accuracy (10-fold Cross-Validation Average): 0.9259648
=====

```

Ridge

```

In [13]: # Select regularization parameter over trainA (50% of training data)
model_lasso <- cv.glmnet(xA, yA, family='binomial',alpha=0)
lambda_min = model_lasso$lambda.min
lambda_1se = model_lasso$lambda.1se

```

lambda.min

```

In [16]: k = 10
n = length(train50B[,1])
fsize = round(n/k)
rmse = rep(0,k)
zolooss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  xB_train = xB[-(((i-1)*fsize+1):(i*fsize)),]
  yB_train = yB[-(((i-1)*fsize+1):(i*fsize))]
  xB_val = xB[((i-1)*fsize+1):(i*fsize),]
  yB_val = yB[((i-1)*fsize+1):(i*fsize)]
  # Fit model on training and make predictions on validation
  model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=0,lambda=lambda_min)
  pred_lo = predict(model_cv, newx = xB_val)
  num_val = length(yB_val)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = yB_val[j]
  }
  # Compute 0-1 loss for each observation
  lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 loss on the val set
  zolooss[i] = mean(lr_loss)
}
xB_train = xB[-(((k-1)*fsize+1):(length(yB))),]
yB_train = yB[-(((k-1)*fsize+1):(length(yB)))]
xB_val = xB[((k-1)*fsize+1):(length(yB)),]
yB_val = yB[((k-1)*fsize+1):(length(yB))]
# Fit model on training and make predictions on validation
model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=0,lambda=lambda_min)
pred_lo = predict(model_cv, newx = xB_val)
num_val = length(yB_val)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = yB_val[j]
}
# Compute 0-1 loss for each observation
lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
# Compute mean 0-1 loss on the val set
zolooss[k] = mean(lr_loss)
test_error_est = mean(zolooss)

cat("=====\n")
cat("GLMNET Ridge Logistic Regression Model with lambda.min\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")

```

```

=====
GLMNET Ridge Logistic Regression Model with lambda.min

Zero-One Loss (10-fold Cross-Validation Average): 0.07603518
Accuracy (10-fold Cross-Validation Average): 0.9239648
=====

```

lambda.1se

```

In [17]: k = 10
n = length(train50B[,1])
fsize = round(n/k)
rmse = rep(0,k)
zolooss = rep(0,k)
for (i in 1:(k-1)){
  # Get train and validation sets
  xB_train = xB[-(((i-1)*fsize+1):(i*fsize)),]
  yB_train = yB[-(((i-1)*fsize+1):(i*fsize))]
  xB_val = xB[((i-1)*fsize+1):(i*fsize),]
  yB_val = yB[((i-1)*fsize+1):(i*fsize)]
  # Fit model on training and make predictions on validation
  model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=0,lambda=lambda_1se)
  pred_lo = predict(model_cv, newx = xB_val)
  num_val = length(yB_val)
  lr_pred = rep(0,num_val)
  actual = rep(0,num_val)
  for (j in 1:num_val){
    if (pred_lo[j]>0){
      lr_pred[j]=1
    }
    actual[j] = yB_val[j]
  }
  # Compute 0-1 loss for each observation
  lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
  # Compute mean 0-1 loss on the val set
  zolooss[i] = mean(lr_loss)
}
xB_train = xB[-(((k-1)*fsize+1):(length(yB))),]
yB_train = yB[-(((k-1)*fsize+1):(length(yB)))]
xB_val = xB[((k-1)*fsize+1):(length(yB)),]
yB_val = yB[((k-1)*fsize+1):(length(yB))]
# Fit model on training and make predictions on validation
model_cv <- glmnet(xB_train, yB_train, family='binomial',alpha=0,lambda=lambda_1se)
pred_lo = predict(model_cv, newx = xB_val)
num_val = length(yB_val)
lr_pred = rep(0,num_val)
actual = rep(0,num_val)
for (j in 1:num_val){
  if (pred_lo[j]>0){
    lr_pred[j]=1
  }
  actual[j] = yB_val[j]
}
# Compute 0-1 loss for each observation
lr_loss = abs(lr_pred-actual) # loss is 0 if NB_pred=actual, 1 otherwise
# Compute mean 0-1 loss on the val set
zolooss[k] = mean(lr_loss)
test_error_est = mean(zolooss)

cat("=====\n")
cat("GLMNET Ridge Logistic Regression Model with lambda.1se\n\n")
cat("Zero-One Loss (10-fold Cross-Validation Average):",test_error_est,"\n")
cat("Accuracy (10-fold Cross-Validation Average):",1-test_error_est,"\n")
cat("=====\n")

```

```

=====
GLMNET Ridge Logistic Regression Model with lambda.1se

Zero-One Loss (10-fold Cross-Validation Average): 0.07703518
Accuracy (10-fold Cross-Validation Average): 0.9229648
=====

```

In []: